



Towards General Learned Representations of TV Shows

Jack Langerman, Zhipeng Fan, Yao Wang

Project Scope



Data




Results



- Validation result (For labels with occurrences > 30)
 - Train on 9536 scenes and validate on 1595 scenes

Model	Precision	Recall	F1
BoW (average)	0.287	0.176	0.198
BoW (tf-idf, norm by # terms)	0.298	0.166	0.190
BoW (tf-idf, norm by sum)	0.366	0.161	0.202
Hierarchical Attention Model	0.237	0.077	0.106
Transformer	0.474	0.277	0.327

- Test result (For labels with occurrences > 30)
 - Test best model (transformer) on test set with 1461 scenes
 - Precision: 0.449 Recall: 0.256 F1: 0.304



Pipelines / Models



[1] Attention is all you need (<https://arxiv.org/abs/1706.03762.pdf>)

